

A mathematical explanation of empirical results in machine learning

Matilde Dolfato

Research Proposal - PhD in Data Science, Università Sapienza di Roma

July 3, 2026

Context

Modern machine learning achieves outstanding predictive performance, but classical statistical learning theory [7, 12] fails to explain many of its successes [3, 5].

Example

Generalization ability of over-parametrized models vs. classical overfitting arguments.

Context

Modern machine learning achieves outstanding predictive performance, but classical statistical learning theory [7, 12] fails to explain many of its successes [3, 5].

Example

Generalization ability of over-parametrized models vs. classical overfitting arguments.

★ We need new mathematical tools to explain successful results ★

Why should we bother?

- ▶ **scientific curiosity**
- ▶ **explainability**
- ▶ **future developments**

“nothing is more practical than a good theory” [11]

State of the art

Current partial progress, on three fronts [5, 9]:

- ▶ **approximation theory**: why deep nets represent complex functions well [6]
- ▶ **optimization theory**: why non-convex objectives are minimized effectively [1]
- ▶ **generalization theory**: why interpolating models still perform well [8]

⇒ no unified theory yet matching the breadth of empirical practice.

Research objective

Research question

Fill the **gap between practice and theory** in deep learning

Research objective

Research question

Fill the **gap between practice and theory** in deep learning

How I intend to do it:

1. **mathematical framework** and hypotheses
2. test hypotheses with **experiments**
3. if discrepancies: **develop the theory** further, then retest
4. if not: run **more complex experiments** to test robustness

Research objective

Research question

Fill the **gap between practice and theory** in deep learning

How I intend to do it:

1. **mathematical framework** and hypotheses
 2. test hypotheses with **experiments**
 3. if discrepancies: **develop the theory** further, then retest
 4. if not: run **more complex experiments** to test robustness
-
- ▶ *definition* $\rightarrow$ *theorem* $\rightarrow$ *proof* [13, 5]
 - ▶ mental pictures (just enough theory)
 - ▶ alternation between theory and practice

Example: the puzzle of generalization


Research question

In the context of gradient-based learning algorithms, **why do over-parametrized models generalize better?**

Example: the puzzle of generalization

Research question

In the context of gradient-based learning algorithms, **why do over-parametrized models generalize better?**




→ Analysis discrete update vs. continuous-time flow (**implicit regularization**)

Double descent - Belkin et al. [4]

Example: the puzzle of generalization

Research question

In the context of gradient-based learning algorithms, **why do over-parametrized models generalize better?**



Double descent - Belkin et al. [4]

→ Analysis discrete update vs. continuous-time flow (**implicit regularization**)

- ▶ mathematical formalization of the problem
- ▶ experiments on gradient-based algorithms dynamics
- ▶ formulate and test hypotheses
- ▶ apply framework to test implicit regularization argument

Thank you

Questions?

References I



Zeyuan Allen-Zhu, Yuanzhi Li, and Zhao Song.

A convergence theory for deep learning via over-parameterization.

In International Conference on Machine Learning, 2019.



David GT Barrett and Benoit Dherin.

Implicit gradient regularization.

arXiv preprint arXiv:2009.11162, 2020.



Mikhail Belkin.

Fit without fear: remarkable mathematical phenomena of deep learning through the prism of interpolation.

Acta Numerica, 30:203–248, 2021.

References II



Mikhail Belkin, Daniel Hsu, Siyuan Ma, and Soumik Mandal.
Reconciling modern machine-learning practice and the classical
bias–variance trade-off.

Proceedings of the National Academy of Sciences,
116(32):15849–15854, 2019.



Julius Berner, Philipp Grohs, Gitta Kutyniok, and Philipp
Petersen.

The modern mathematics of deep learning.

arXiv preprint arXiv:2105.04026, 78:3, 2021.



Helmut Bölcskei, Philipp Grohs, Gitta Kutyniok, and Philipp
Petersen.

Optimal approximation with sparsely connected deep neural
networks.

SIAM Journal on Mathematics of Data Science, 2019.

References III



Trevor Hastie, Robert Tibshirani, Jerome Friedman, et al.
The elements of statistical learning, 2009.



Behnam Neyshabur, Srinadh Bhojanapalli, David McAllester,
and Nathan Srebro.

Towards understanding the role of over-parametrization in
generalization of neural networks.

In International Conference on Learning Representations, 2019.



Tomaso Poggio, Andrzej Banburski, and Qianli Liao.

Theoretical issues in deep networks: Approximation,
optimization and generalization.

Proceedings of the National Academy of Sciences, 2020.

References IV



Samuel L Smith, Benoit Dherin, David GT Barrett, and Soham De.

On the origin of implicit regularization in stochastic gradient descent.

arXiv preprint arXiv:2101.12176, 2021.



Vladimir Vapnik.

The nature of statistical learning theory.

Springer science & business media, 2013.



Martin J Wainwright.

High-dimensional statistics: A non-asymptotic viewpoint, volume 48.

Cambridge university press, 2019.

References V



E Weinan.

A mathematical perspective of machine learning.

In *International Congress of Mathematicians*, pages 914–954.

European Mathematical Society-EMS-Publishing House
GmbH, 2023.

Appendix I

Proposition [Comparing discrete and continuous trajectories]
[2, 10]

For a step size $\alpha \in \mathbb{R}_{>0}$, $t \in \mathbb{R}_{>0}$

$$\theta_{t+\alpha} = \theta_t - \alpha \nabla \ell(\theta_t) \quad \text{and} \quad \frac{d\theta}{dt} = -\nabla \ell(\theta(t))$$

Up to order α , gradient descent corresponds to gradient flow of:

$$\ell_{discr=contin}(\theta) = \ell(\theta) + \frac{\alpha}{4} \|\nabla \ell(\theta)\|^2$$

Proof:

Consider the usual (discrete-time) update rule:

$$\theta_{\tau+1} = \theta_{\tau} - \alpha \nabla \ell(\theta_{\tau})$$

for $\tau \in \mathbb{N}$, $\alpha \in \mathbb{R}_{>0}$ step size.

Appendix II

Let $t = \alpha\tau$ and consider the finite-difference

$$\theta_{t+\alpha} - \theta_t = -\alpha \nabla \ell(\theta_t).$$

Now, taking the limit we obtain the continuous-time update rule (where we omit the dependence of θ on t to ease notation)

$$\lim_{\alpha \rightarrow 0} \frac{\theta_{t+\alpha} - \theta_t}{\alpha} = \frac{d\theta}{dt}.$$

Then,

$$\frac{d\theta}{dt} = -\nabla \ell(\theta). \quad (1)$$

Allowing

$$\frac{d\theta}{dt} = -\nabla \ell(\theta) + \alpha g(\theta),$$

we want to find the correction term g such that the continuous and discrete trajectories agree up to order α . Let us consider the

Appendix III

function $\theta(t)$ and inspect its Taylor expansion starting from $\theta_0 = \theta(0)$, for a jump $\Delta t = \alpha$ (1 iteration in discrete sense):

$$\begin{aligned}\theta(t) &\approx \theta_0 + \alpha \frac{d\theta}{dt} + \frac{1}{2}\alpha^2 \frac{d^2\theta}{dt^2} \\ &= \theta_0 - \alpha \nabla \ell(\theta) + \alpha^2 g(\theta) + \frac{1}{2}\alpha^2 \frac{d}{dt} \left(-\nabla \ell(\theta) + \alpha g(\theta) \right) \\ &\approx \theta_0 - \alpha \nabla \ell(\theta) + \alpha^2 \left(g(\theta) + \frac{1}{2} \left(-\nabla^2 \ell(\theta) \frac{d\theta}{dt} \right) \right) \\ &= \theta_0 - \alpha \nabla \ell(\theta) + \alpha^2 \left(g(\theta) + \frac{1}{2} \left(\nabla^2 \ell(\theta) \left(\nabla \ell(\theta) + \alpha g(\theta) \right) \right) \right)\end{aligned}$$

where the terms in pink are of 3rd order in α , hence disregarded. Then, we impose $\theta(t) = \theta_{\tau+1}$, obtaining

$$g(\theta) = -\frac{1}{2} \nabla^2 \ell(\theta) \nabla \ell(\theta).$$

Appendix IV

Hence, the following continuous-time update rule agrees with the discrete-time update up to order α

$$\frac{d\theta}{dt} = -\nabla\ell(\theta) - \frac{1}{2}\alpha\nabla^2\ell(\theta)\nabla\ell(\theta). \quad (2)$$

Lastly, we notice that this is the same as applying the usual continuous-time update rule to the modified loss

$$\tilde{\ell}(\theta) = \ell(\theta) + \frac{1}{4}\alpha\nabla\ell(\theta)^t\nabla\ell(\theta) = \ell(\theta) + \frac{1}{4}\alpha\|\nabla\ell(\theta)\|^2 \quad (3)$$